# GenAI Nantes

## "Comment sur-optimiser Claude Code pour shipper les yeux fermés"

**Cédric Chariere Fiedler**

CEO @ siliceum | Ingénieur fiabilité & performance

**18 Juin** 2026 à 19h00

3 Rue Lavoisier, 44100 Nantes

chez onepoint.

Sii
techtown
lonestone
RCA
externatic
zenika
_icilundi

# GenAI Nantes

## "Postmortem: comment mon agent de debug a pris le contrôle de ma relation client"

**Cédric Chariere Fiedler**
CEO @ Crédit Agricole | Testeur

**18 Juin** 2026 à 19h00

3 Rue Lavoisier, 44100 Nantes

chez onepoint.

SII techtown lonestone RCA externatic zenika _icilundi

# GenAI Nantes

🎤 15 événements / an
🏴‍☠️ 1 hackathon (Shift)
👷 2 workshops
❤️ 1 communauté de 1400p*

* 14.000p selon le syndicat des llamas 🦙

# Merci aux partenaires ! ❤️

💛 Gold

💛 Silver

lonestone

techtown

zenika

externatic

RCA ÉDITEUR DE LOGICIELS POUR LES EXPERTS-COMPTABLES

# Qui est développeur ? 🪪

# Les autres, vous faites quoi dans la vie ? 🪪

# Qui utilise Claude Code ? 🤖

# Qui utilise Codex ? 🤖

# Qui utilise Antigravity ? 🤖

# Qui utilis(ait) Github Copilot ? 🤖

# Autre ? 🤖

# Qui recrute dans la GenAI ? 🚀

# Qui cherche un job dans la GenAI ? 🏋️

# Qui souhaite s’associer dans la GenAI ? 👨‍❤️‍👨

Tu es disponible pour des missions **GenAI** ?

# Schedule

👀 1- News (infrastructures)

🏋 2- Claude Code

🍻 3- Enjoy

# News !


**Softbank investit 75 G€ dans des datacenters IA Français**

45 G€ pour 3.1 GW (2031)

3 G€ pour 2 GW

# Décryptons

# Datacenter stack

Greenfield -> rack
- Shell
- Power
- Network
- Cooling

# Datacenter stack

Compute

- Servers
- GPU
- Storage
- Switch

# Datacenter stack

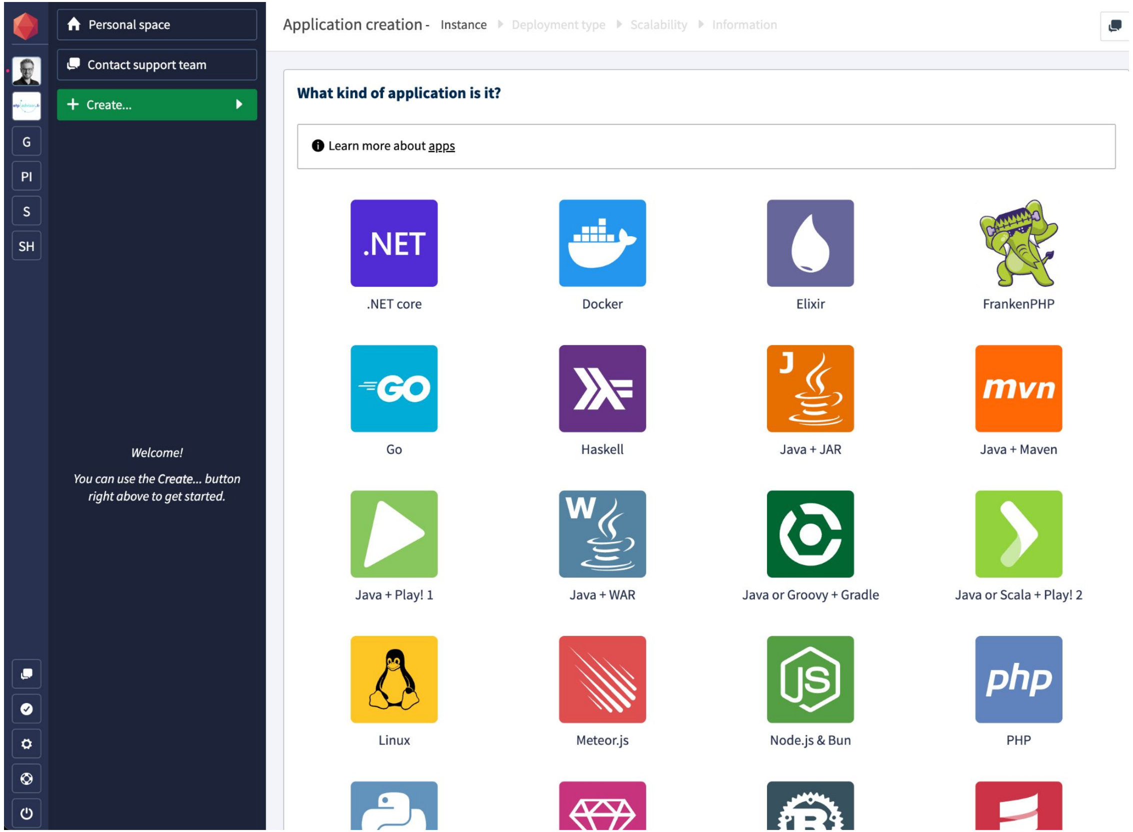


Cloud / Software

- IaaS
- PaaS
- Applications
- Databases

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| Google / Alphabet | Build or lease | Buys | GCP |
| Amazon / AWS | Build or lease | Buys | AWS |
| Microsoft | Build or lease | Buys | Azure |
| OVH / Scaleway | Build | Buys | IaaS |
| Clever Cloud | | Buys or lease | PaaS |
| Heroku | | | PaaS |

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| Equinix | Build | | |
| Telehouse | Build | | |
| Data4 | Build | | |
| SoftBank (France) | Build | | |

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| CoreWeave | | Buys NVIDIA | Inference endpoint |
| Nebius | | Buys NVIDIA | Inference endpoint |

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| NVIDIA | | Own chips | WIP |
| AMD | | Own chips | |
| Cerebras | | Own chips | Inference endpoint |

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| OpenAI | WIP (Stargate) | WIP (Stargate) | Inference endpoint |
| Anthropic | | WIP | Inference endpoint |
| xAI | Build (Colossus) | Buys NVIDIA | Inference endpoint |
| Mistral (FR) | WIP | WIP | Inference endpoint |

# Datacenter stack

| Actor | Shell | Hardware | Software / Cloud |
|---|---|---|---|
| Google / Alphabet | Build or lease | Buys | GCP |
| Amazon / AWS | Build or lease | Buys | AWS |
| Microsoft | Build or lease | Buys | Azure |
| OVH / Scaleway | Build | Buys | IaaS |
| Clever Cloud | | Buys or lease | PaaS |
| Heroku | | | PaaS |
| | | | |
| NVIDIA | | Own chips | WIP |
| AMD | | Own chips | |
| Cerebras | | Own chips | Inference endpoint |
| | | | |
| Equinix | Build | | |
| Telehouse | Build | | |
| Data4 | Build | | |
| SoftBank (France) | Build | | |
| | | | |
| CoreWeave | | Buys NVIDIA | Inference endpoint |
| Nebius | | Buys NVIDIA | Inference endpoint |
| | | | |
| OpenAI | WIP (Stargate) | WIP (Stargate) | Inference endpoint |
| Anthropic | | WIP | Inference endpoint |
| xAI | Build (Colossus) | Buys NVIDIA | Inference endpoint |
| Mistral (FR) | WIP | WIP | Inference endpoint |

# AI datacenters CapEx per GW

| Actor | Bucket | € per GW |
|---|---|---|
| Shell | DC construction (shell, land, fit-out) | ~2.5 G€ |
| | Cooling (CDUs, chillers, towers, liquid loops) | ~2.5 G€ |
| | Electrical (substation, transformers, switchgear, UPS, gensets) | ~3.5 G€ |
| | Network (NICs, switches, optics, NVLink backplane) | ~5 G€ |
| | Other (spares, integration, commissioning) | ~1.5 G€ |
| Hardware | Servers excluding-GPU (CPU, NVSwitch, board, power delivery) | ~10 G€ |
| | GPU | ~15 G€ |
| Software | Software | ~0.5 G€ |
| | **Total up-front CapEx** | **~40.5 G€** |

# AI datacenters CapEx per GW

| Actor | Bucket | € per GW |
|---|---|---|
| **Shell** | DC construction (shell, land, fit-out) | ~15 G€ |
| | Cooling (CDUs, chillers, towers, liquid loops) | |
| | Electrical (substation, transformers, switchgear, UPS, gensets) | |
| | Network (NICs, switches, optics, NVLink backplane) | |
| | Other (spares, integration, commissioning) | |
| **Hardware** | Servers excluding-GPU (CPU, NVSwitch, board, power delivery) | ~25 G€ |
| | GPU | |
| **Software** | Software | ~0.5 G€ |
| | **Total up-front CapEx** | **~40.5 G€** |

# AI datacenters CapEx per GW

| Actor | Bucket | € per GW |
|---|---|---|
| Shell | DC construction (shell, land, fit-out) | ~15 G€ |
| | Cooling (CDUs, chillers, towers, liquid loops) | |
| | Electrical (substation, transformers, switchgear, UPS, gensets) | |
| | Network (NICs, switches, optics, NVLink backplane) | |
| | Other (spares, integration, commissioning) | |
| Hardware | Servers excluding-GPU (CPU, NVSwitch, board, power delivery) | ~25 G€ |
| | GPU | |
| Software | Software | ~0.5 G€ |
| | **Total up-front CapEx** | **~40.5 G€** |

Softbank -> 3 GW / 45 G€

# Investments Oil & Gas vs AI

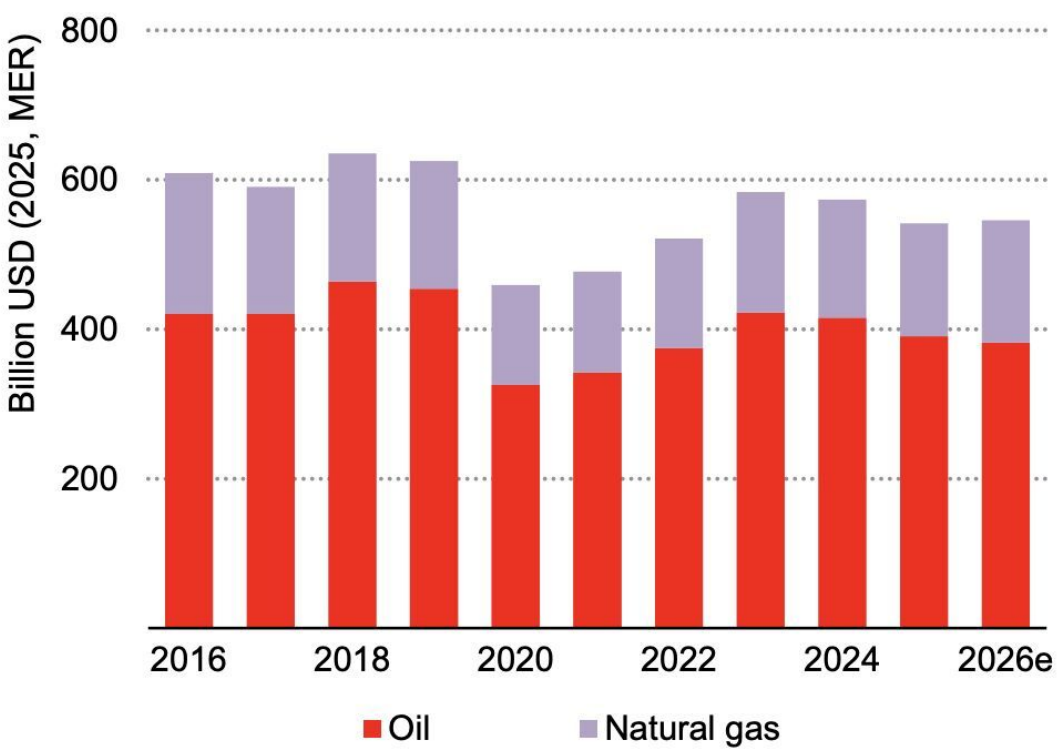


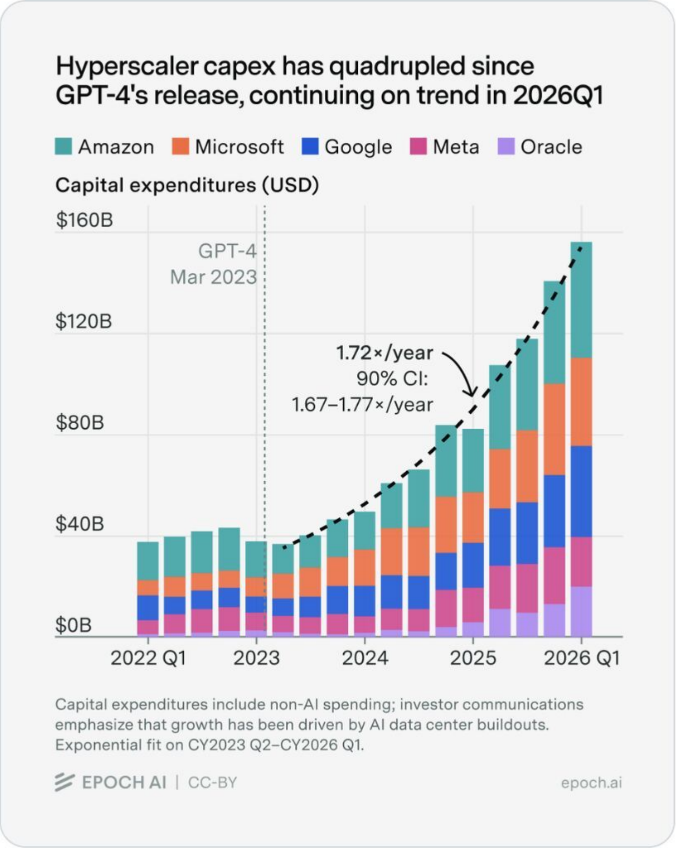

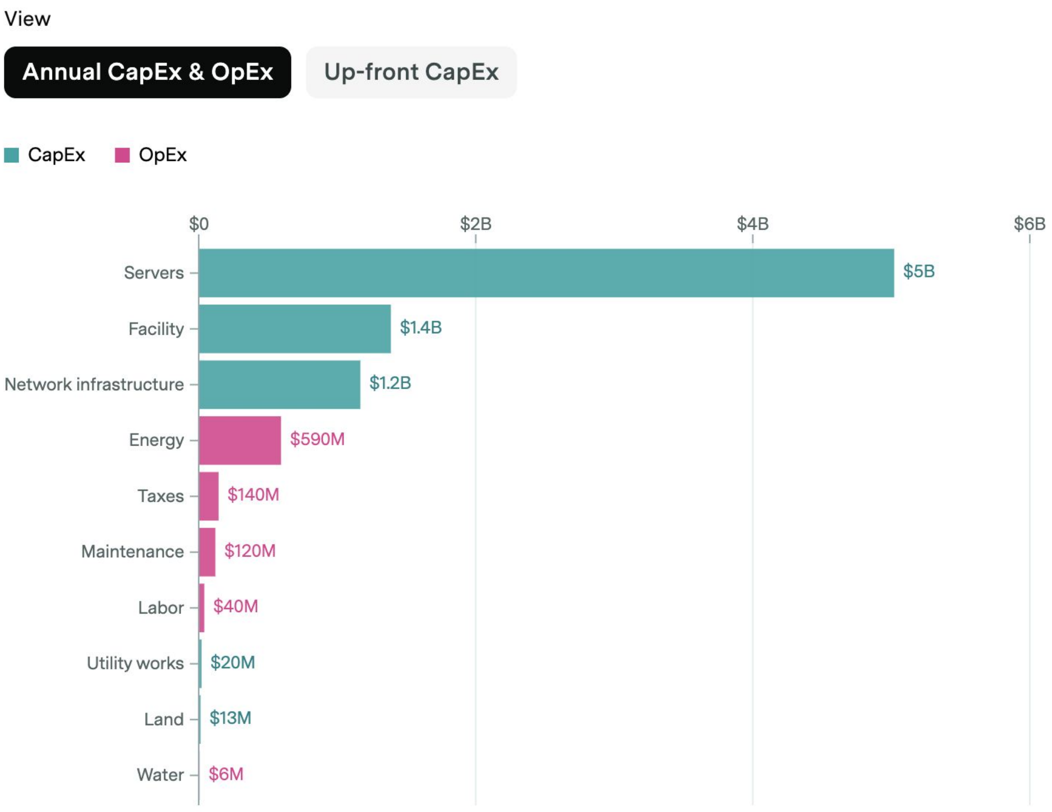


Based on a stylized model of a US hyperscaler AI data center that has 1 GW of nameplate capacity for the IT equipment. CapEx is converted to annual costs over each asset's lifespan, using the cost of capital as a discount rate. "Facility" includes the building shell, mechanical systems, and electrical equipment.



Download graph

# TCO - amortized

# Make France rich again

| Actor | Net income |
|---|---|
| Shell | 10% |
| TSMC | 45% |
| GPUs (Nvidia) | 60% |
| Inference (Mistral - projected) | 50% |

# Make France rich again

SESTERCE

VSORA®

iliad

MISTRAL
AI_

# Make France rich again


Arthur Mensch: “1.000 G€ / year of AI revenues in Europe”

# Power generation

OK:

- Nuclear power
- Gas turbine

KO:

- Solar panels
- Wind turbine

# Power generation

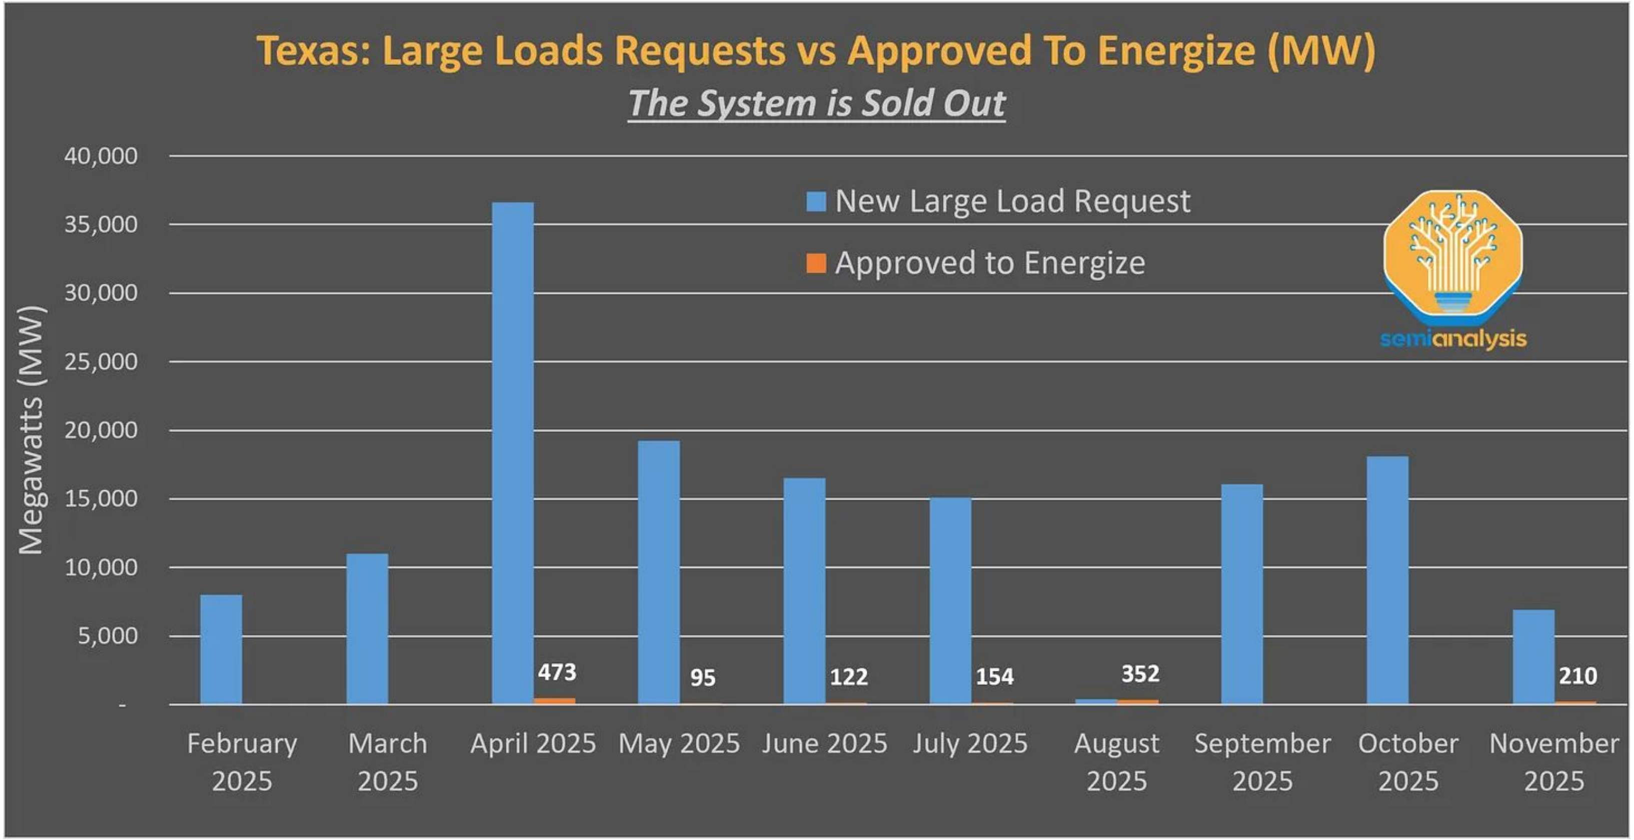

# Power generation

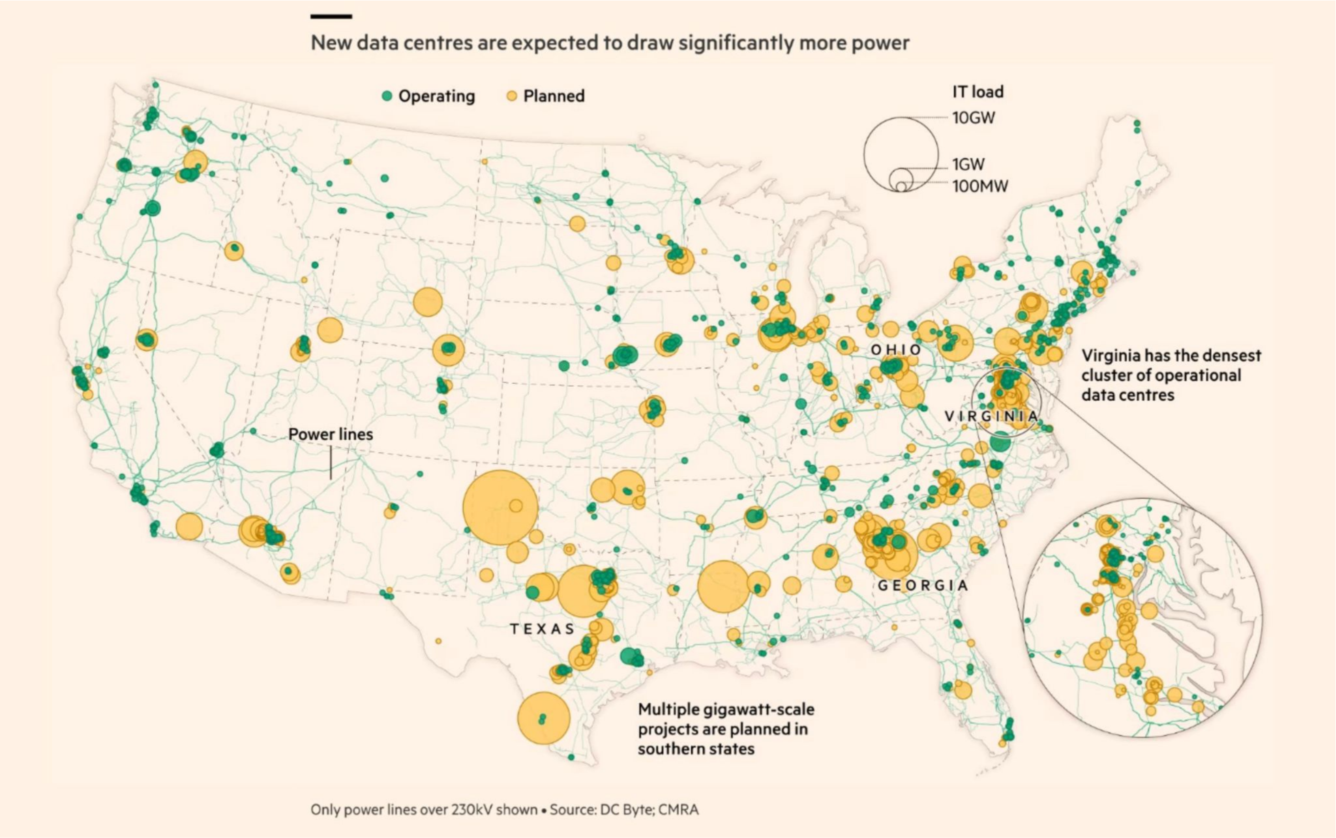

# Power generation

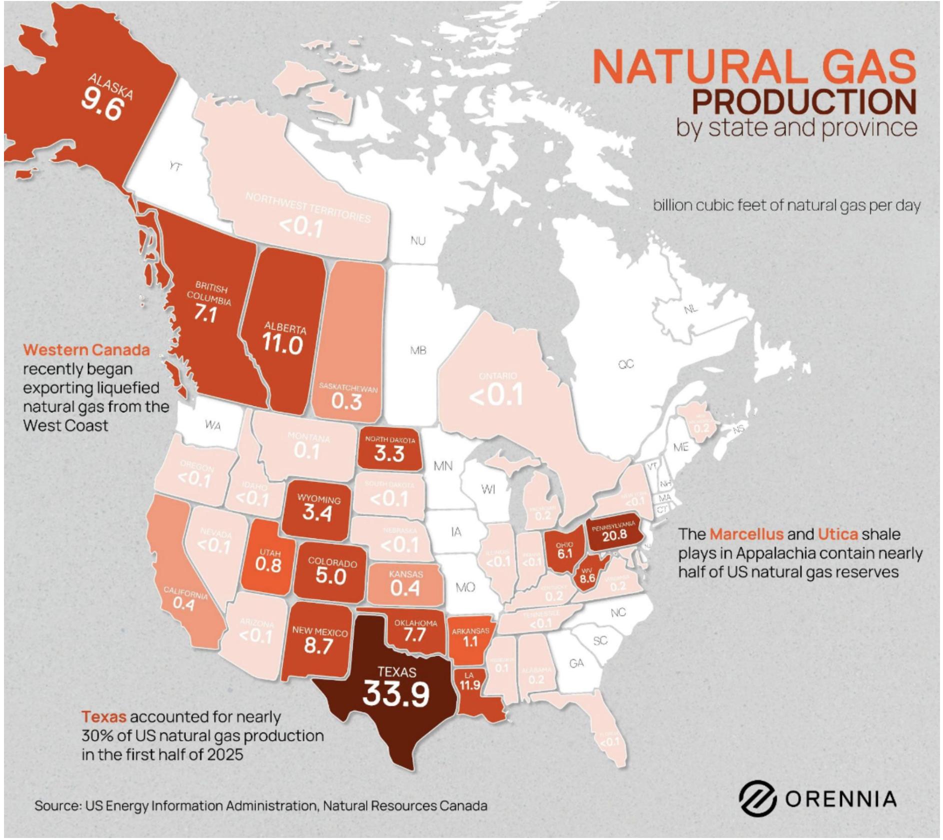

# Power generation (Colossus 2)

# Power generation (Colossus 2)

# Power generation (Colossus 2)

# Power generation (Colossus 2)

# Power generation (Colossus 2)

# Power generation (Colossus 2)

## Global natural gas turbine orders, 2024-2025

Last updated 16 Apr 2026

Download chart ↓

Cite Share

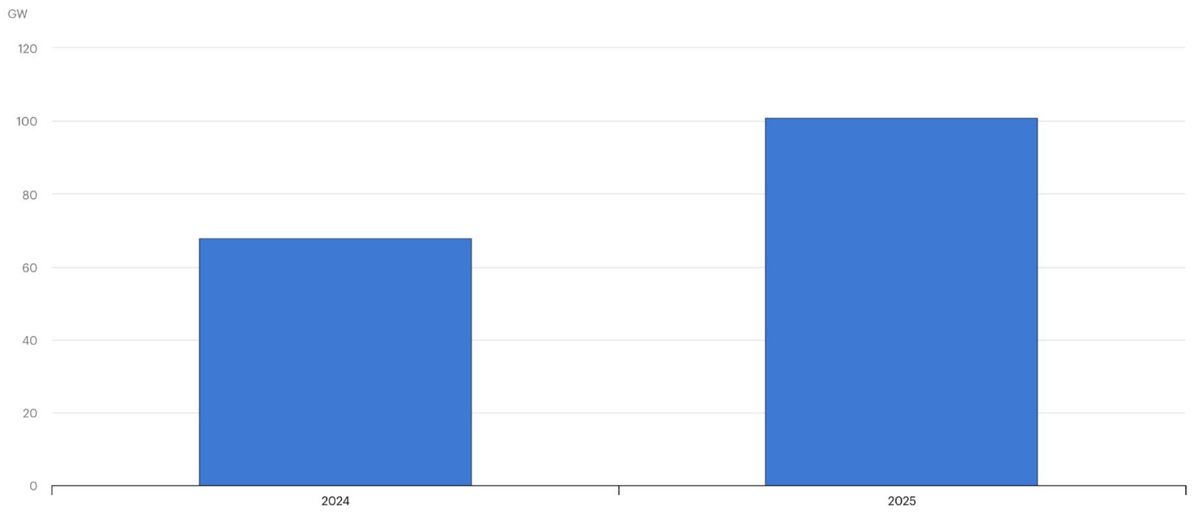

# Power generation

NUCLEAR / NUCLEAR POWER

## Supercarrier USS Gerald R. Ford To Act As Floating Nuclear Power Plant For Facilities On Land

The Pentagon is exploring ways to keep the power on at critical bases after attacks or natural disasters, and there's a history of ships acting in this role.

JOSEPH TREVITHICK / PUBLISHED MAY 24, 2026 4:43 PM EDT / 193 / ADD TWZ

USN

# Rapport RTE “Futurs énergétiques 2050”

Production actuelle: 550 TWh
Consommation actuelle: 450 TWh
Excédents: 100 TWh

Electrification transports: 75 TWh
Electrification industrie: 50 TWh
Electrification bâtiments: 50 TWh

Datacenters IA en projet: 25 TWh

# “Effet falaise” du parc nucléaire

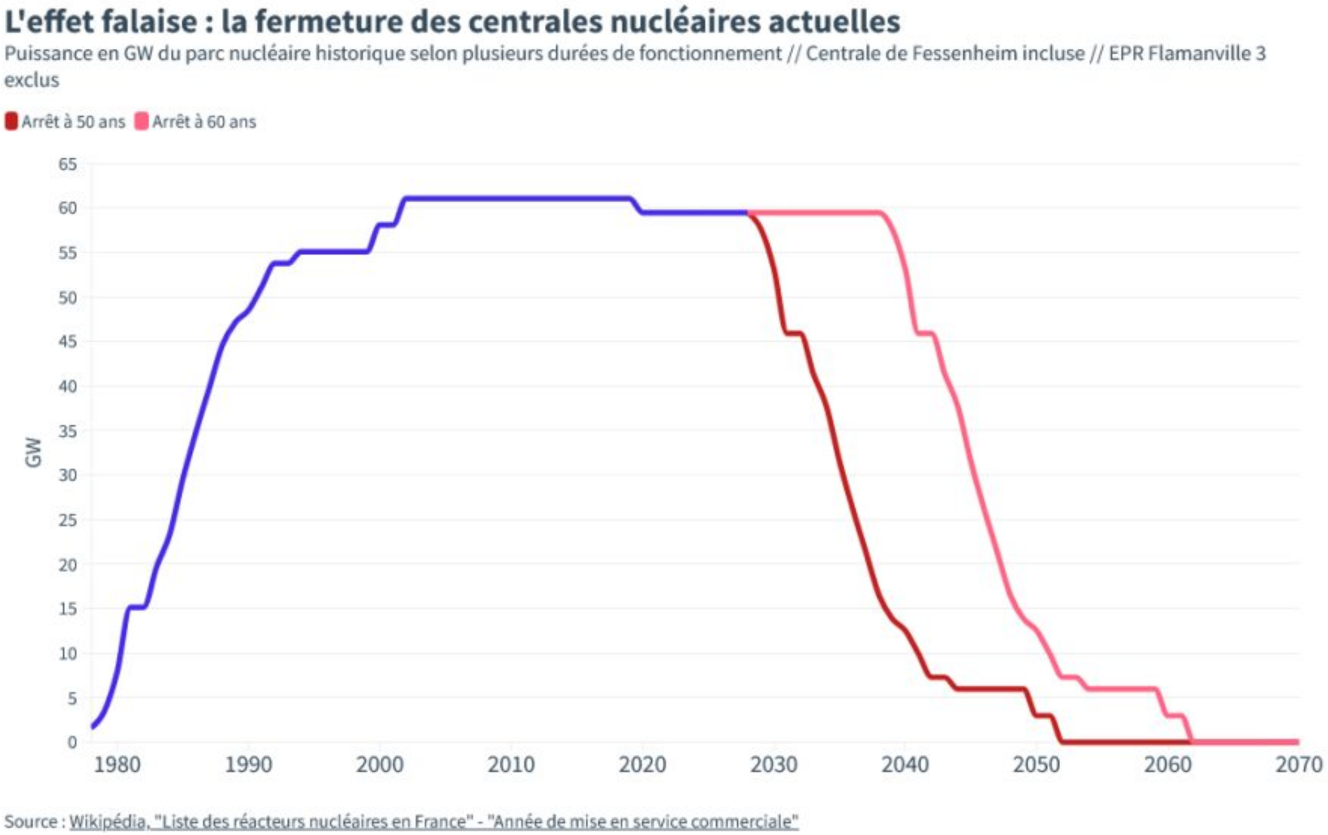

# RTE policy: “premier connecté, premier servi”

Available according to RTE: 10 GW

Projects:
AI datacenters: 10 GW
Industry: 10 GW

# RTE policy: “premier connecté, premier servi”

* SoftBank, Hauts-de-France: 5 GW, 3.1 GW of it by 2031
* Campus IA (MGX, Nvidia, Mistral, Bpifrance), south of Paris: 1.4 GW
* Data4 / Brookfield: ~1.5 GW
* AION (Iliad-Scaleway, Orange, EDF, Ardian, Bull): ~1 GW
* Google, Amazon, Microsoft expansions: ~1 GW combined

# RTE policy: “premier connecté, premier servi”

Closing in 2027
(maybe, maybe not)

# Greenhouse gas emissions

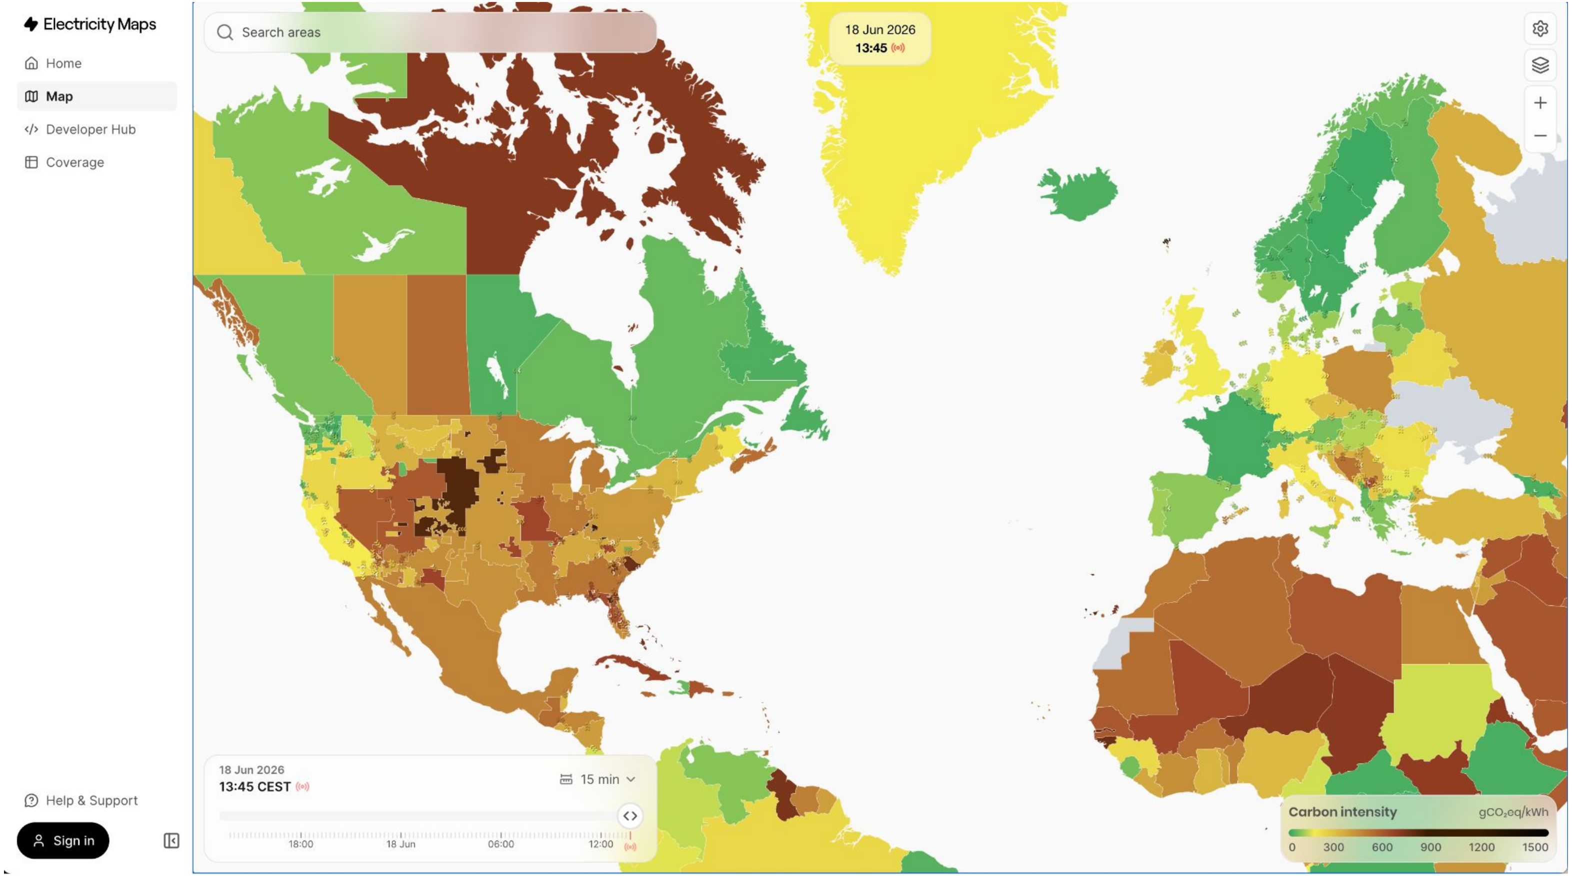

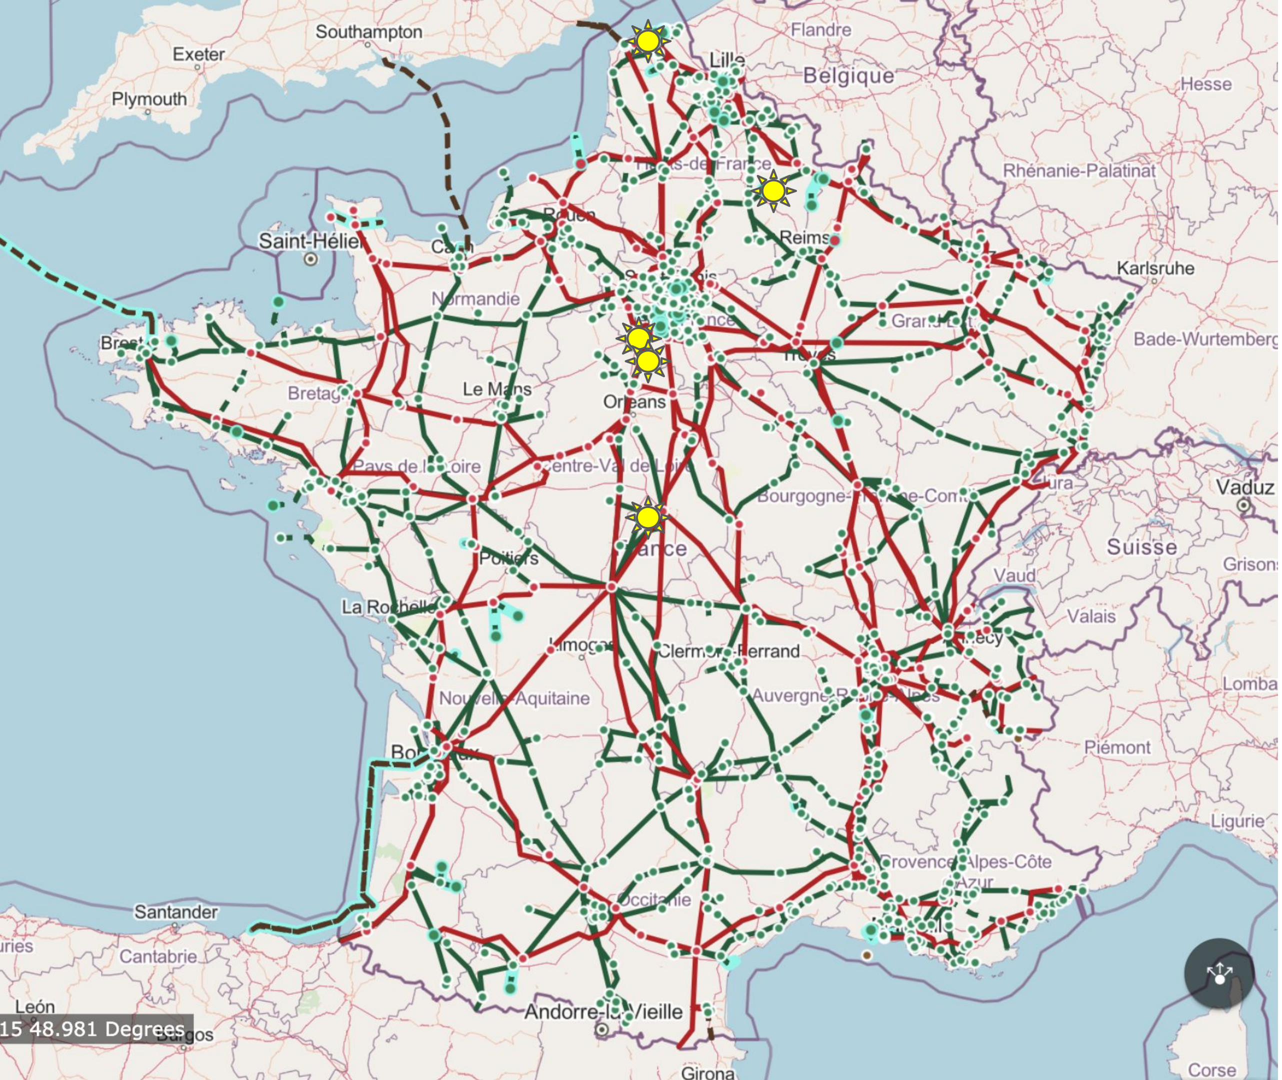


# French AI datacenters map

# Softbank datacenter project (3GW)

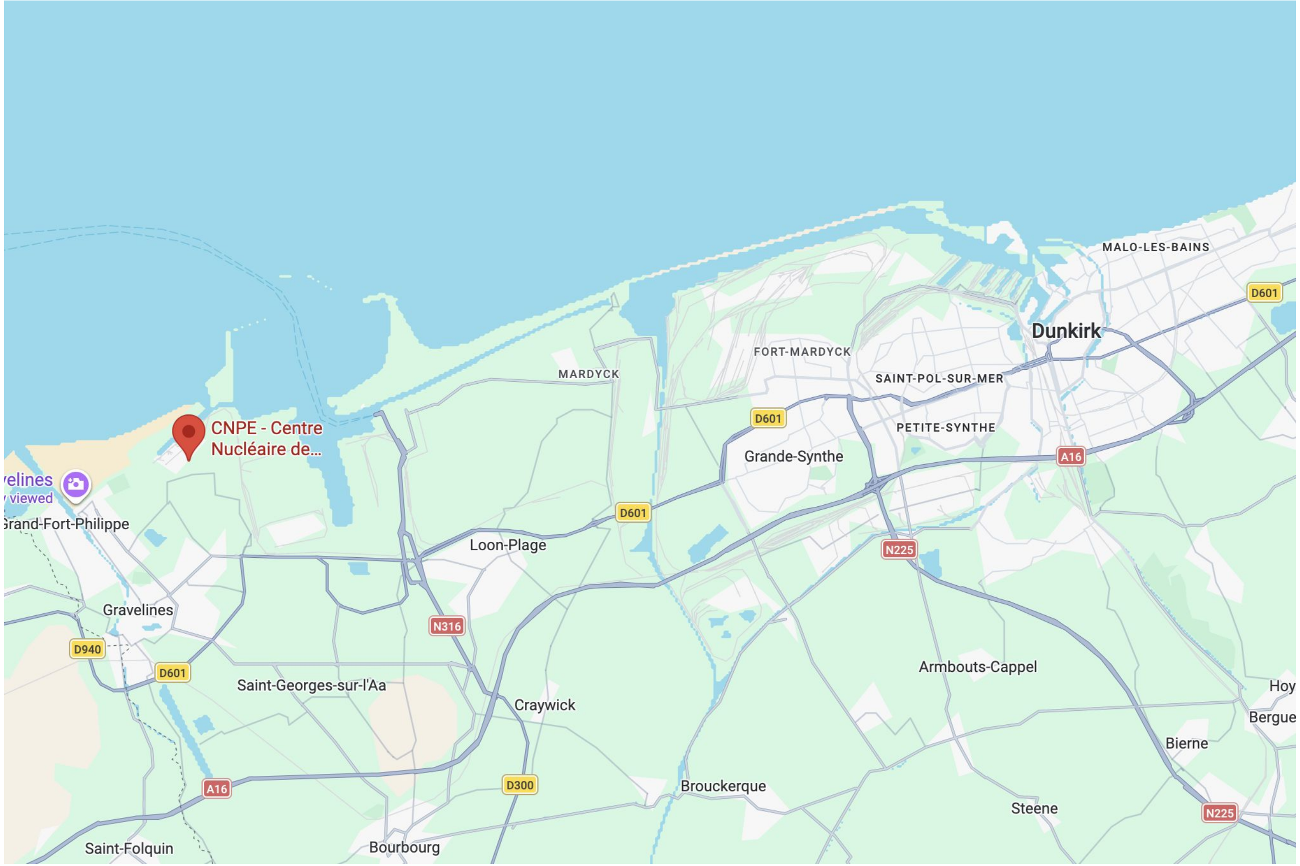

# SpaceX (Colossus 1) x Anthropic

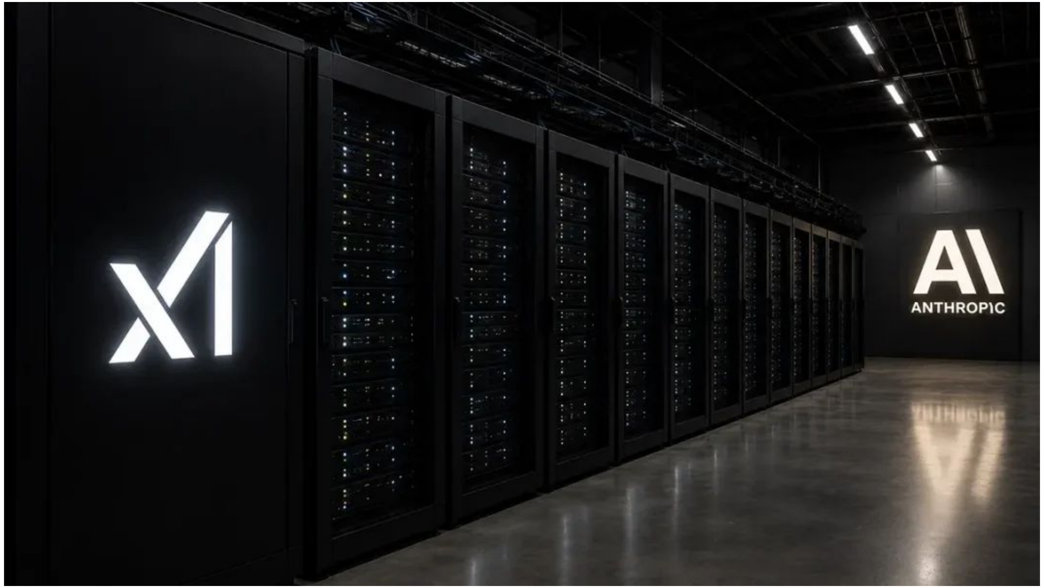


Colossus 1:

- training originally
- inference now (Anthropic)

# SpaceX (Colossus 1) x Anthropic

**Training**:

- Nvidia GPUs
- 100k GPUs for a single task
- Limited by the slowest GPU
- InfiniBand / East-West traffic
- Flat load (100%)

**Inference**:

- Nvidia GPUs or ASIC
- 1 to 72 GPUs for a single task
- Ethernet / North-South traffic
- KV cache on SSD + offloading
- Variable load

# xAI (Colossus 1) x Anthropic

Colossus 1 mostly on H100

Training moved to Colossus 2

Drop in Grok usage

Colossus 1 is on Hopper
Colossus 2 is on Blackwell
Training require a uniform platform

# Datacenters in Space

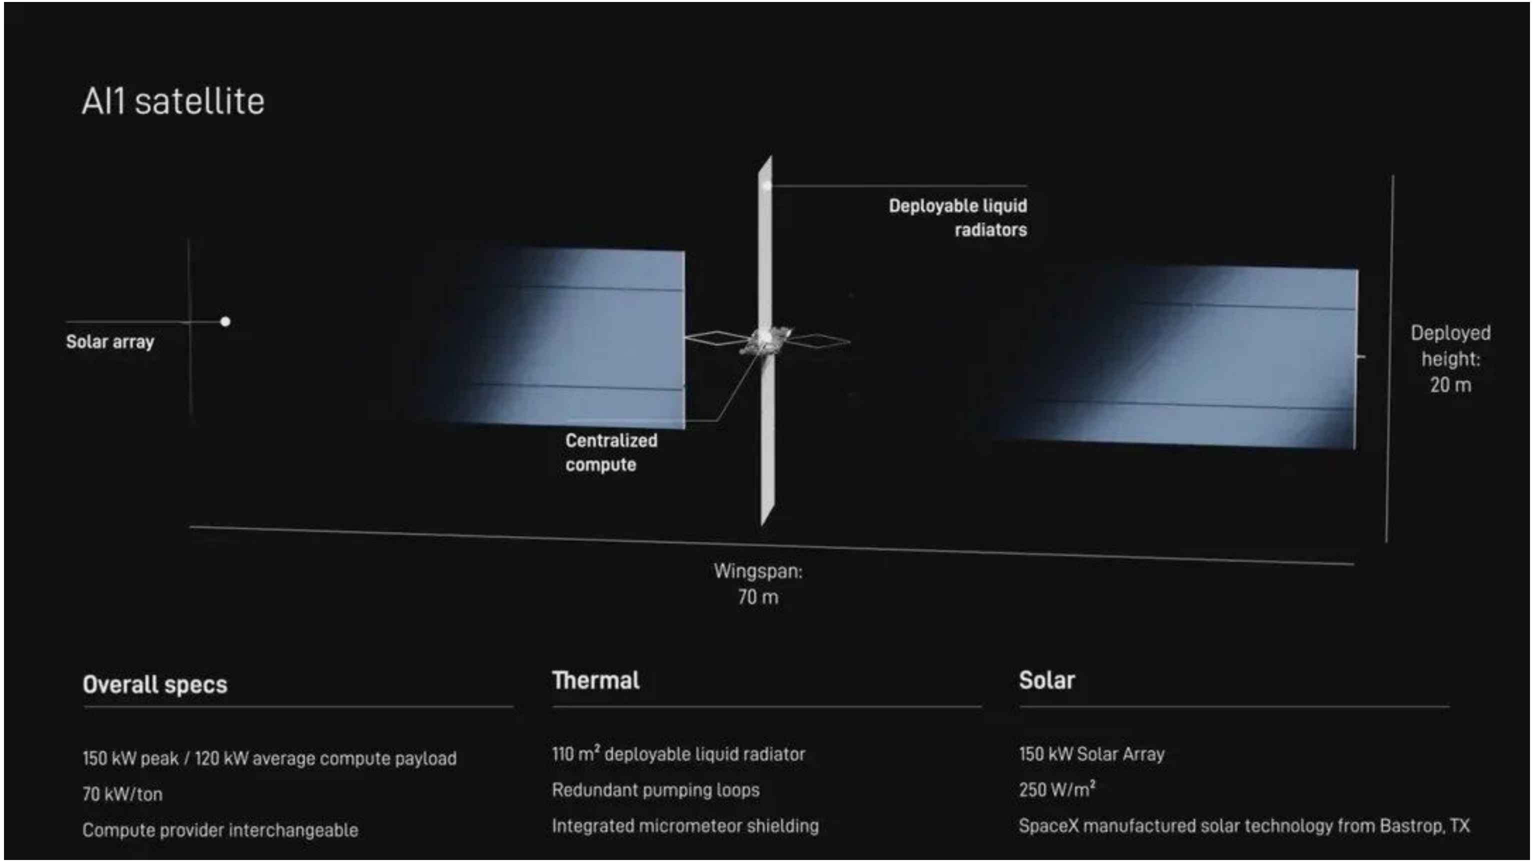

# Big datacenters vs constellation

View

Annual CapEx & OpEx | Up-front CapEx

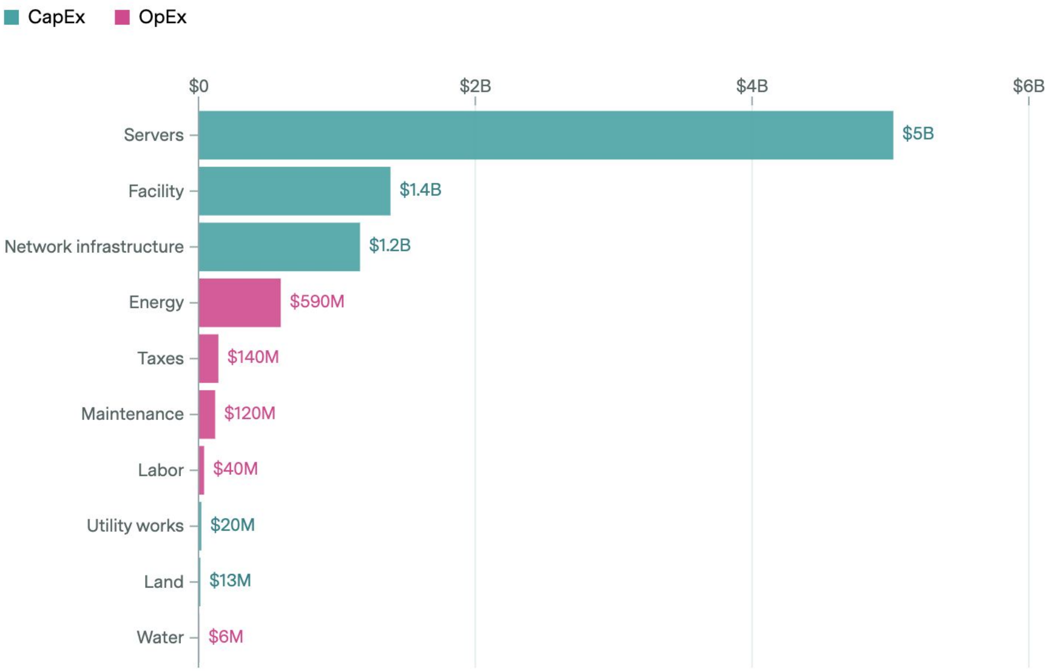


Based on a stylized model of a US hyperscaler AI data center that has 1 GW of nameplate capacity for the IT equipment. CapEx is converted to annual costs over each asset's lifespan, using the cost of capital as a discount rate. "Facility" includes the building shell, mechanical systems, and electrical equipment.



Download graph

# On earth

# Datacenters in Space (according to Elon)

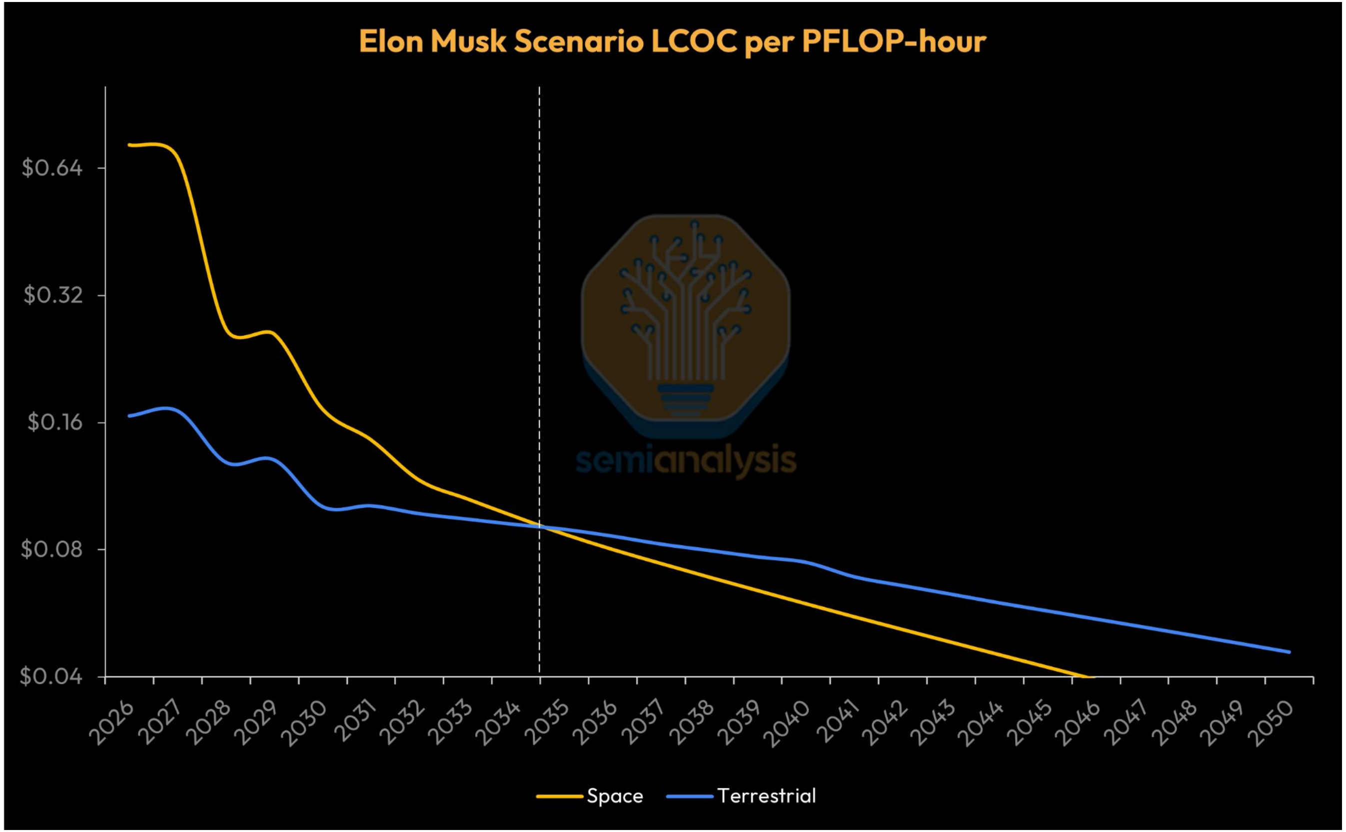

# Datacenters in Space (according to SemiAnalytics)

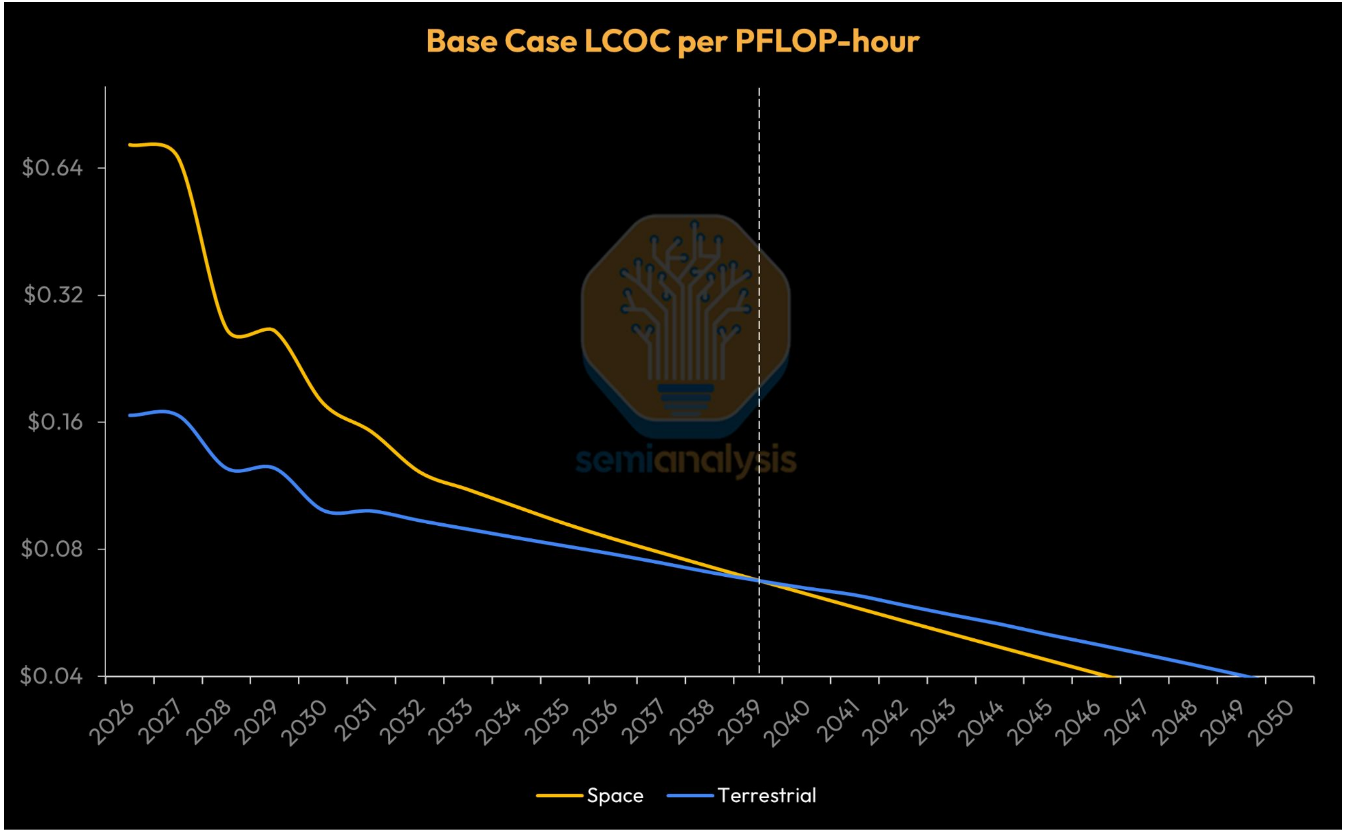

# TCO (according to SemiAnalytics)

| AI Cloud Total Cost of Ownership Summary: Space vs Terrestrial (Extended Prices) | | | |
|---|---|---|---|
| | Unit | B300 1200W (Orbital) | B300 1200W |
| Site | | Space | Terrestrial |
| Customer Profile | | Hyperscaler | Hyperscaler |
| Year | | 2026 | 2026 |
| Program Power | W | 30,565 | 30,565 |
| Power per GPU | W | 1,910 | 1,910 |
| GPUs | GPUs | 16 | 16 |
| GPU per Server | GPUs | 8 | 8 |
| IT Cluster Capital Cost | USD | $980,882 | $986,158 |
| Datacenter Capital Cost[1] | USD | $3,086,332 | $382,061 |
| **Total Program Capital Cost** | **USD** | **$4,067,215** | **$1,368,219** |
| Weighted Average Cost of Capital | % | 15.0% | 10.3% |
| Datacenter Useful Life in Years | Years | 5 | 15 |
| Monthly IT Cluster Capital Cost of Ownership | USD | $23,335 | $21,099 |
| Monthly Datacenter Capital Cost of Ownership[1] | USD | $73,424 | $4,176 |
| Monthly Operating Cost of Ownership | USD | $4,167 | $2,449 |
| **Total Monthly Cost of Ownership** | **USD** | **$100,925** | **$27,724** |

*1. Includes launch cost.*

# Cooling

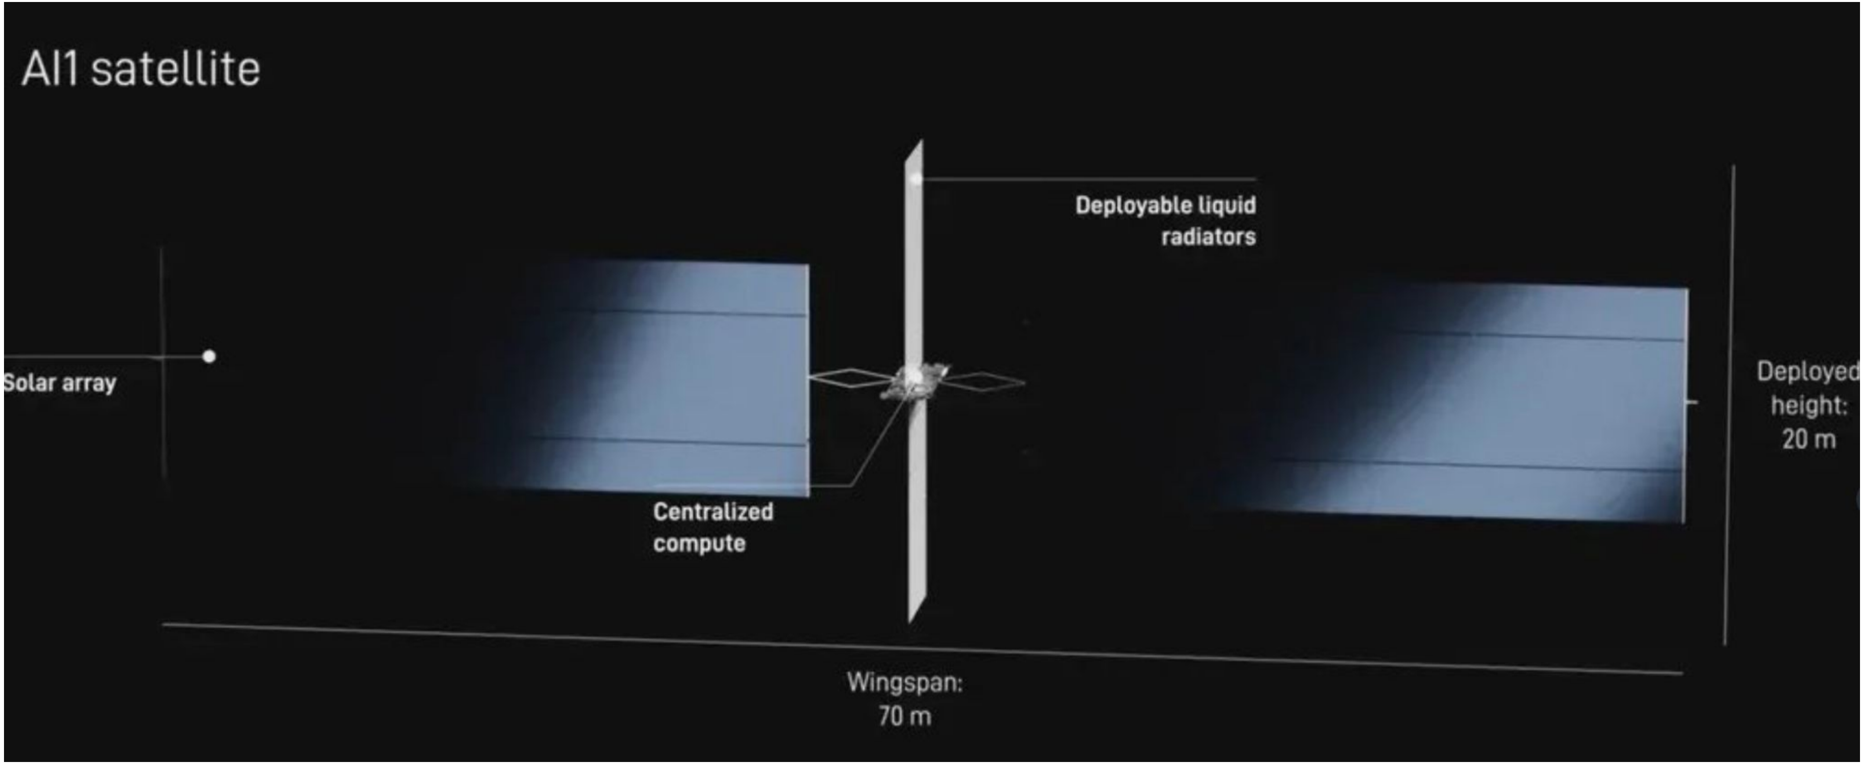
AI1 satellite
Deployable liquid radiators
Solar array
Deployed height: 20 m
Centralized compute
Wingspan: 70 m

# Datacenters in Space (orbit)

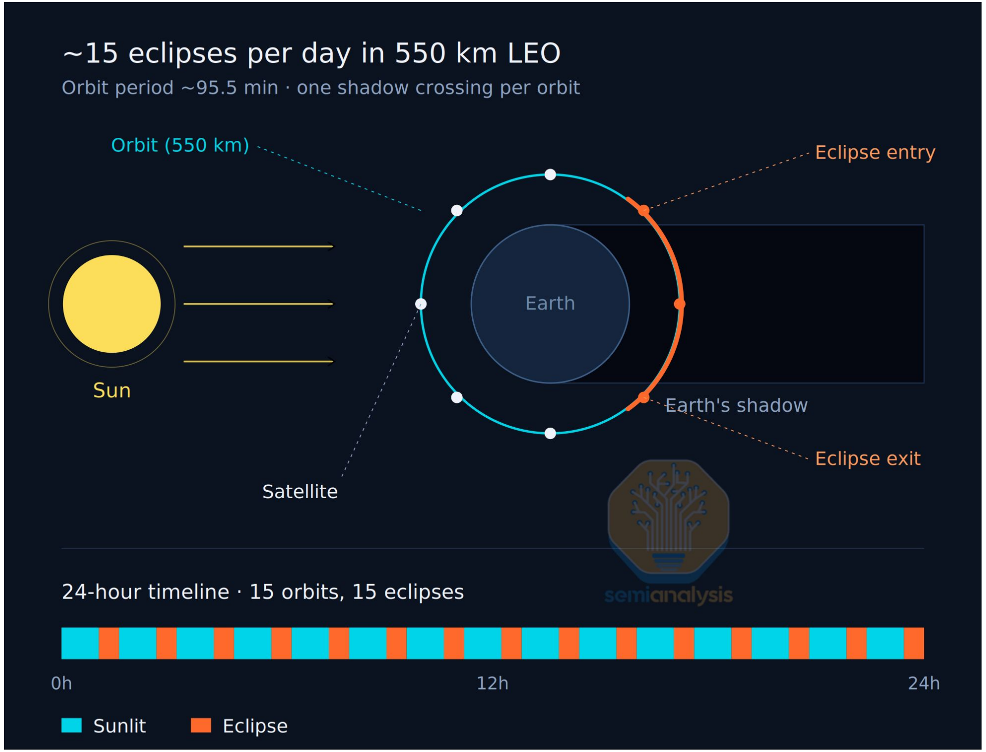


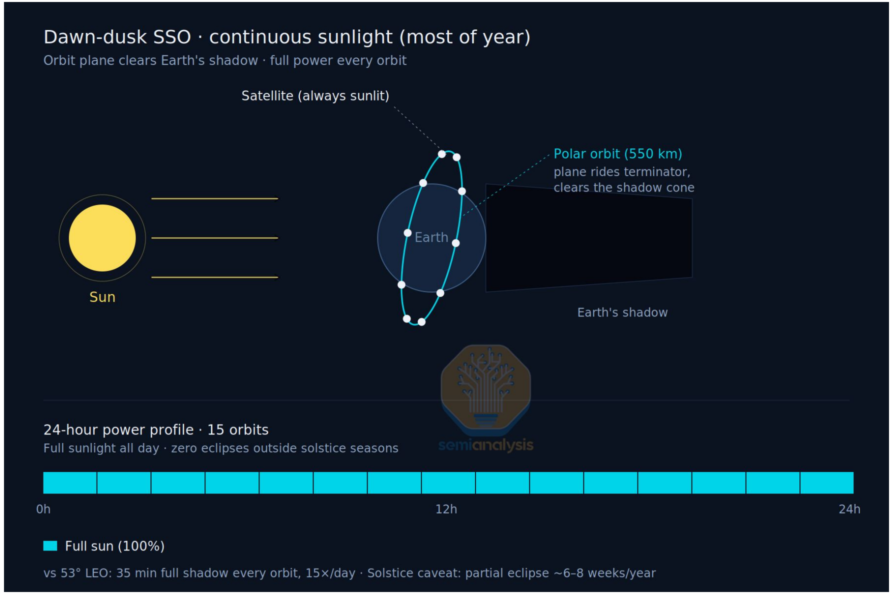

# Network

LEO
Lots of ground stations (5-7 min)
Large oceans and empty zones

High bandwidth (for KV-cache)

30-80ms one-way

https://newsletter.semianalysis.com/p/to-boldl
y-go-the-case-for-space-datacenters

Announcements

# Claude Fable 5 and Claude Mythos 5

Jun 9, 2026

PRESIDENTIAL ACTIONS

# PROMOTING ADVANCED ARTIFICIAL INTELLIGENCE INNOVATION AND SECURITY

Executive Orders | June 2, 2026

By the authority vested in me as President by the Constitution and the laws of the United States of America, it is hereby ordered:

Section 1. Purpose. The United States continues to lead the world in Artificial Intelligence (AI) because of the enormous talent and innovation of our AI industry, and because we refuse to stifle this innovation with overly burdensome regulation. My Administration has unleashed tremendous technological growth and economic investment in AI by slashing the

# Identity verification on Claude

Updated today



Being responsible with powerful technology starts with knowing who is using it. Identity verification helps us prevent abuse, enforce our usage policies, and comply with legal obligations.

We are rolling out identity verification for a few use cases, and you might see a verification prompt when accessing certain capabilities, as part of our routine platform integrity checks, or other safety and compliance measures.

We only use your verification data to confirm who you are and not for any other purposes.

## How are we verifying?

We selected Persona Identities as our verification partner based on the strength of their technology, privacy controls, and security safeguards. Follow the steps below to complete your identity verification process.

## What you'll need

Before you start, have these ready:

- **A valid government-issued photo ID**: the physical document, in hand
- **A phone or a computer with a camera**: you may be asked to take a live selfie with your phone, or your webcam
- **A few minutes**: verification typically takes under five minutes

MISTRAL AI_

PRODUCT ANNOUNCEMENT | 14 JUNE 2026

# Mistral introduces Le Chaton Fat

A 30T+ parameter frontier model built for long-horizon reasoning, coding, and agentic work.

| 30T+ | Sparse MoE | 1M | Native | 256k | 50+ |
|---|---|---|---|---|---|
| total parameters | architecture | token context window | multimodal | output | languages |

## State-of-the-art performance across leading benchmarks

| Benchmark | Le Chaton Fat | Claude Fable 5 | Claude Opus 4.8 | GPT-5.5 |
|---|---|---|---|---|
| MMLU-Pro (Pass@1) | 90.6 | 88.9 | 84.2 | 82.6 |
| GPQA Diamond (Pass@1) | 82.4 | 80.8 | 74.8 | 73.1 |
| SWE-bench Verified (Resolved) | 76.8 | 73.4 | 66.1 | 67.4 |
| HumanEval (Pass@1) | 96.1 | 94.7 | 90.3 | 89.7 |

Benchmark comparisons against publicly available launch materials for Claude Fable 5 as of June 14, 2026, plus internal evaluation harnesses.

Available in Le Chat, API, and enterprise deployments.

mistral.ai

# GenAI Nantes

## “Comment sur-optimiser Claude Code pour shipper les yeux fermés”

**Cédric Chariere Fiedler**

CEO @ siliceum | Ingénieur fiabilité & performance

**18 Juin** 2026 à 19h00

4 Rue Voltaire, 44000 Nantes

chez onepoint.

SII techtown lonestone RCA

externatic zenika _icilundi

**Slides** dispo sur:
https://github.com/genai-nantes-meetup/meetups

# GenAI
## Nantes

## Workshop: “Midjourney de A à Z.”

**Simon Timssale**
Consultant et formateur GenAI

**Yannis Sulfourt**
Consultant et formateur GenAI

**4 Mars** 2026 à 19h00

2 Pl. Louis Daubenton, 44100 Nantes

chez {EPITECH} LEARN DIFFERENT*

sii techtown lonestone RCA
externatic zenika


On vous ❤
1
2
Restez pour l'after-work !